

# Modeling and Machine Learning in R: tidymodels

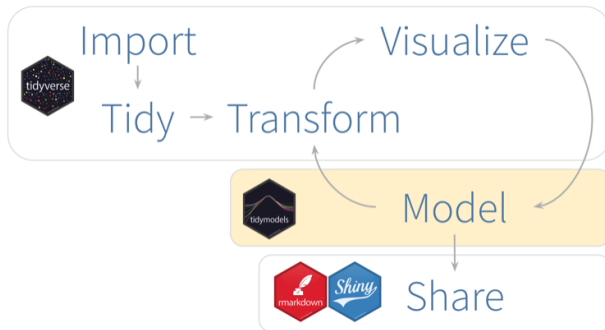
TIAGO SOUZA

# Acknowledgements

The practical example in the slides is based on a blog post **Tutorial on tidymodels for Machine Learning** by Hansjörg Plieninger.

<https://hansjoerg.me/2020/02/09/tidymodels-for-machine-learning/>

# tidymodels within the R Universe




- **tidymodels** is to MODELING what the **tidyverse** is to DATA WRANGLING;
- **tidymodels** has a modular approach: specific, smaller packages are designed to work hand in hand.

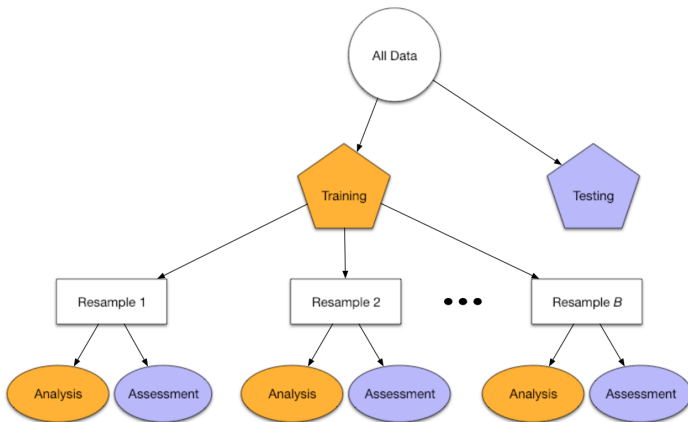
# tidymodels' main packages



# Goal: predict diamond prices



# How are we going to do it?



# What tools do we have?

Pre-Process → Train → Validate



# Separating Testing and Training Data



- *rsample* contains a set of functions to create different types of resamples and corresponding classes for their analysis:
  - Traditional resampling techniques for estimating the sampling distribution of a statistic and;
  - Estimating model performance using a holdout set.



# Separating Testing and Training Data

```
dia_split <- initial_split(diamonds, pro = .1, strata = price)

dia_train <- training(dia_split)
dia_test  <- testing(dia_split)

dia_vfold <- vfold_cv(dia_train, v = 3, repeats = 1, strata = price)
print(dia_vfold)
```

```
## # 3-fold cross-validation using stratification
## # A tibble: 3 x 2
##   splits          id
##   <list>         <chr>
## 1 <split [3594/1799]> Fold1
## 2 <split [3595/1798]> Fold2
## 3 <split [3597/1796]> Fold3
```

# Data Pre-Processing and Feature Engineering



- *recipes* is a method for creating and pre-processing design matrices used for modeling or visualization;
- Idea: define a blueprint that can be used to sequentially define the encodings and pre-processing of the data;
- It is used to prepare a data set (for modeling) using different 'step\_\*()' functions;
- The 'recipe()' takes a formula and a data set, and then the different steps are added.

# Data Pre-Processing and Feature Engineering

```
dia_rec <-  
  recipe(price ~ ., data = dia_train) %>%  
    step_log(all_outcomes()) %>%  
    step_normalize(all_predictors(), -all_nominal()) %>%  
    step_dummy(all_nominal()) %>%  
    step_poly(carat, degree = 2)  
  
prep(dia_rec)
```

```
## Recipe  
##  
## Inputs:  
##  
##      role #variables  
## outcome      1  
## predictor      9  
##  
## Training data contained 5393 data points and no missing data.  
##  
## Operations:  
##  
## Log transformation on price [trained]  
## Centering and scaling for carat, depth, table, x, y, z [trained]  
## Dummy variables from cut, color, clarity [trained]  
## Orthogonal polynomials on carat [trained]
```

# Data Pre-Processing and Feature Engineering

- Calling 'prep()' on a recipe applies all steps;
- Call 'juice()' to extract the transformed data set;
- Call 'bake()' on a new data set.

```
dia_juiced <- juice(prepare(dia_rec))  
names(dia_juiced)
```

```
## [1] "depth"      "table"      "x"          "y"          "z"  
## [6] "price"      "cut_1"      "cut_2"      "cut_3"      "cut_4"  
## [11] "color_1"    "color_2"    "color_3"    "color_4"    "color_5"  
## [16] "color_6"    "clarity_1"  "clarity_2"  "clarity_3"  "clarity_4"  
## [21] "clarity_5"  "clarity_6"  "clarity_7"  "carat_poly_1" "carat_poly_2"
```

# Defining and Fitting Models



- The goal is to provide a tidy, unified interface to models that can be used to try a range of models without getting bogged down in the syntactical minutiae of the underlying packages;
- Has wrappers around many popular machine learning algorithms, and you can fit then using a unified interface.

# Defining and Fitting Models

- ① Function specific to each algorithm;
- ② 'set\_mode()' (regression or classification);
- ③ 'set\_engine()' back-end/engine/implementation

```
lm_model <-  
  linear_reg() %>%  
  set_mode("regression") %>%  
  set_engine("lm")  
  
print(lm_model)
```

```
## Linear Regression Model Specification (regression)  
##  
## Computational engine: lm
```

# Defining and Fitting Models

- Random Forest: 'ranger' or 'randomForest'?
- How to handle their different interfaces?

```
rand_forest(mtry = 3, trees = 500, min_n = 5) %>%  
  set_mode("regression") %>%  
  set_engine("ranger", importance = "impurity_corrected")
```

```
## Random Forest Model Specification (regression)  
##  
## Main Arguments:  
##   mtry = 3  
##   trees = 500  
##   min_n = 5  
##  
## Engine-Specific Arguments:  
##   importance = impurity_corrected  
##  
## Computational engine: ranger
```

# Defining and Fitting Models

- This example, with a formula. You can also set 'x' and 'y'.

```
lm_fit1 <- fit(lm_model, price ~ ., dia_juiced)
lm_fit1
```

```
## parsnip model object
##
## Fit time: 13ms
##
## Call:
## stats::lm(formula = price ~ ., data = data)
##
## Coefficients:
## (Intercept)          depth          table             x             y
##    7.729257    0.009178    0.003533    0.042671    0.295996
##           z          cut_1          cut_2          cut_3          cut_4
##    0.038552    0.081400   -0.001890    0.004473   -0.004180
##    color_1    color_2    color_3    color_4    color_5
##   -0.442545   -0.086210   -0.010589    0.022553   -0.009546
##    color_6    clarity_1    clarity_2    clarity_3    clarity_4
##    0.001549    0.836278   -0.215930    0.101644   -0.050261
##    clarity_5    clarity_6    clarity_7    carat_poly_1    carat_poly_2
##    0.005939   -0.007926    0.028804    49.458215   -16.835516
```



# Summarizing Fitted Models



- Takes the messy output of built-in function in R, such as 'lm', 'nls', and turns them into tidy tibbles;
- From **tidyverse**.

# Summarizing Fitted Models

- 'glance()' reports information about the entire model;
- 'tidy()' summarizes information about model components.

```
glance(lm_fit1$fit)
```

```
## # A tibble: 1 x 12
##   r.squared adj.r.squared sigma statistic p.value    df logLik    AIC    BIC
##   <dbl>      <dbl> <dbl>    <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
## 1    0.978        0.978 0.149   10064.    0     24  2629. -5206. -5035.
## # ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

```
tidy(lm_fit1) %>%
  arrange(desc(abs(statistic))) %>%
  print()
```

```
## # A tibble: 25 x 5
##   term          estimate std.error statistic  p.value
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)    7.73     0.00410  1885.    0
## 2 clarity_1      0.836     0.0125   66.7    0
## 3 color_1       -0.443     0.00706  -62.7    0
## 4 carat_poly_2  -16.8      0.284    -59.3    0
## 5 carat_poly_1   49.5       1.26     39.3    4.70e-297
## 6 clarity_2     -0.216     0.0117   -18.5    9.45e- 74
## 7 color_2      -0.0862     0.00646  -13.3    5.65e- 40
## 8 y             0.296     0.0269   11.0     8.93e- 28
## 9 clarity_3      0.102     0.0100   10.1     5.79e- 24
## 10 cut_1         0.0814     0.00948   8.58    1.19e- 17
## # ... with 15 more rows
```

# Summarizing Fitted Models


- 'augment()' is used to get model predictions, residuals, etc.

```
lm_predicted <- augment(lm_fit1$fit, data = dia_juiced) %>%  
  rowid_to_column()  
print(select(lm_predicted, rowid, price, .fitted:.std.resid))
```

```
## # A tibble: 5,393 x 8  
##   rowid price .fitted .resid   .hat .sigma   .cooksd .std.resid  
##   <int> <dbl>   <dbl> <dbl>   <dbl> <dbl>   <dbl>   <dbl>  
## 1     1  5.86    6.18 -0.320 0.00378 0.149 0.000704 -2.15  
## 2     2  5.87    6.04 -0.174 0.00444 0.149 0.000244 -1.17  
## 3     3  6.00    6.26 -0.264 0.00439 0.149 0.000557 -1.78  
## 4     4  6.00    6.45 -0.452 0.00750 0.149 0.00280 -3.05  
## 5     5  6.31    6.48 -0.162 0.00288 0.149 0.000138 -1.09  
## 6     6  6.32    6.56 -0.239 0.00493 0.149 0.000511 -1.61  
## 7     7  6.32    6.13  0.191 0.00426 0.149 0.000281  1.28  
## 8     8  6.32    6.44 -0.123 0.00233 0.149 0.0000642 -0.828  
## 9     9  6.32    6.45 -0.124 0.00253 0.149 0.0000711 -0.837  
## 10    10  6.32    6.50 -0.176 0.00398 0.149 0.000224 -1.18  
## # ... with 5,383 more rows
```

# Visualizing Results

```
ggplot(lm_predicted, aes(.fitted, price)) +  
  geom_point(alpha = .2) +  
  ggrepel::geom_label_repel(aes(label = rowid),  
    data = filter(lm_predicted, abs(.resid) > 2)) +  
  labs(title = "Actual vs. Predicted Price of Diamonds")
```



# Evaluating Model Performance



# Evaluating Model Performance

- Use 'rsample', 'parsnip' and 'yardstick' for cross-validation (3).

```
print(dia_vfold)

## # 3-fold cross-validation using stratification
## # A tibble: 3 x 2
##   splits          id
##   <list>         <chr>
## 1 <split [3594/1799]> Fold1
## 2 <split [3595/1798]> Fold2
## 3 <split [3597/1796]> Fold3
```

- Extract analysis/training and assessment/testing data.

```
lm_fit2 <- mutate(dia_vfold,
                  df_ana = map(splits, analysis),
                  df_ass = map(splits, assessment))
print(lm_fit2)

## # 3-fold cross-validation using stratification
## # A tibble: 3 x 4
##   splits          id  df_ana          df_ass
##   <list>         <chr> <list>         <list>
## 1 <split [3594/1799]> Fold1 <tibble [3,594 x 10]> <tibble [1,799 x 10]>
## 2 <split [3595/1798]> Fold2 <tibble [3,595 x 10]> <tibble [1,798 x 10]>
## 3 <split [3597/1796]> Fold3 <tibble [3,597 x 10]> <tibble [1,796 x 10]>
```

# Evaluating Model Performance

- Prepare data / fit model / predict.

```
lm_fit2 <-  
  lm_fit2 %>%  
    # prep, juice, bake  
    mutate(  
      recipe = map(df_ana, ~prep(dia_rec, training = .x)),  
      df_ana = map(recipe, juice),  
      df_ass = map2(recipe, df_ass, ~bake(.x, new_data = .y))  
    ) %>%  
    #fit  
    mutate(  
      model_fit = map(df_ana, ~fit(lm_model, price ~ ., data = .x))  
    ) %>%  
    # predict  
    mutate(  
      model_pred = map2(model_fit, df_ass, ~predict(.x, new_data = .y))  
    )  
  
print(select(lm_fit2, id, recipe:model_pred))
```

```
## # A tibble: 3 x 4  
##   id      recipe  model_fit model_pred  
##   <chr> <list>    <list>    <list>  
## 1 Fold1 <recipe> <fit[+]> <tibble [1,799 x 1]>  
## 2 Fold2 <recipe> <fit[+]> <tibble [1,798 x 1]>  
## 3 Fold3 <recipe> <fit[+]> <tibble [1,796 x 1]>
```

# Evaluating Model Performance

- Select original and predicted values.

```
lm_preds <-  
  lm_fit2 %>%  
  mutate(res = map2(df_ass, model_pred, ~data.frame(price = .x$price,  
                                                    .pred = .y$.pred))  
  ) %>%  
  select(id, res) %>%  
  tidyr::unnest(res) %>%  
  group_by(id)  
  
print(lm_preds)
```

```
## # A tibble: 5,393 x 3  
## # Groups:   id [3]  
##   id    price .pred  
##   <chr> <dbl> <dbl>  
## 1 Fold1  6.00  6.28  
## 2 Fold1  6.31  6.48  
## 3 Fold1  6.32  6.57  
## 4 Fold1  6.32  6.45  
## 5 Fold1  6.32  6.56  
## 6 Fold1  6.32  6.45  
## 7 Fold1  6.32  6.23  
## 8 Fold1  6.33  6.48  
## 9 Fold1  6.33  6.41  
## 10 Fold1 6.33  6.46  
## # ... with 5,383 more rows
```




# Evaluating Model Performance

- 'metrics()' has default measures for numeric and categorical outcomes (numeric - 'rmse', 'rsq', 'mae');
- You can choose other if you'd like with 'metric\_set()'.

```
print(metrics(lm_preds, truth = price, estimate = .pred))
```

```
## # A tibble: 9 x 4
##   id      .metric .estimator .estimate
##   <chr> <chr>      <chr>      <dbl>
## 1 Fold1 rmse      standard    0.141
## 2 Fold2 rmse      standard    0.237
## 3 Fold3 rmse      standard    0.147
## 4 Fold1 rsq       standard    0.981
## 5 Fold2 rsq       standard    0.945
## 6 Fold3 rsq       standard    0.978
## 7 Fold1 mae       standard    0.114
## 8 Fold2 mae       standard    0.110
## 9 Fold3 mae       standard    0.114
```

# Tuning Model Parameters



- 'tune' wants to facilitate hyper-parameter tuning for the **tidymodels** packages;
- 'dials' contains tools to create and manage values of tuning parameters;
- Let's tune the 'mtry' and 'degree' parameters.

# Tuning Model Parameters

- Preparing a 'parsnip' Model for tuning.

```
rf_model <-  
  rand_forest(mtry = tune()) %>%  
  set_mode("regression") %>%  
  set_engine("ranger")  
  
print(parameters(rf_model))
```

```
## Collection of 1 parameters for tuning  
##  
## identifier type      object  
##           mtry mtry nparam[?]  
##  
## Model parameters needing finalization:  
##   # Randomly Selected Predictors ('mtry')  
##  
## See '?dials::finalize' or '?dials::update.parameters' for more information.
```


# Tuning Model Parameters

- Preparing Data for Tuning: 'recipes';
- Tune the degree of the polynomial for the variable 'carat'.

```
dia_rec2 <-  
  recipe(price ~ ., data = dia_train) %>%  
    step_log(all_outcomes()) %>%  
    step_normalize(all_predictors(), -all_nominal()) %>%  
    step_dummy(all_nominal()) %>%  
    step_poly(carat, degree = tune())  
  
dia_rec2 %>%  
  parameters() %>%  
  pull("object") %>%  
  print()
```

```
## [[1]]  
## Polynomial Degree (quantitative)  
## Range: [1, 3]
```

# Combine Everything



- Object that can bundle together pre-processing, modeling and post-processing requests;
- The recipe prepping and model fitting can be executed using a single call to 'fit()'.

# Combine Everything

```
rf_wflow <-  
  workflow() %>%  
    add_model(rf_model) %>%  
    add_recipe(dia_rec2)  
  
print(rf_wflow)
```

```
## == Workflow =====  
## Preprocessor: Recipe  
## Model: rand_forest()  
##  
## -- Preprocessor -----  
## 4 Recipe Steps  
##  
## * step_log()  
## * step_normalize()  
## * step_dummy()  
## * step_poly()  
##  
## -- Model -----  
## Random Forest Model Specification (regression)  
##  
## Main Arguments:  
##   mtry = tune()  
##  
## Computational engine: ranger
```

# Tuning Parameters

- Update the parameters in the workflow;
- Cross-validation for tuning: select the best combination of hyper-parameters.

```
rf_param <-  
  rf_wflow %>%  
  parameters() %>%  
  update(mtry = mtry(range = c(3L, 5L)),  
         degree = degree_int(range = c(2L, 4L)))  
  
print(rf_param$object)
```

```
## [[1]]  
## # Randomly Selected Predictors (quantitative)  
## Range: [3, 5]  
##  
## [[2]]  
## Polynomial Degree (quantitative)  
## Range: [2, 4]
```

# Tuning Parameters

```
rf_grid <- grid_regular(rf_param)


print(rf_grid)
```

```
## # A tibble: 9 x 2
##   mtry degree
##   <int> <int>
## 1     3     2
## 2     4     2
## 3     5     2
## 4     3     3
## 5     4     3
## 6     5     3
## 7     3     4
## 8     4     4
## 9     5     4
```



# Tuning Parameters

```
rf_search <- tune_grid(rf_wflow,  
  grid = rf_grid,  
  resamples = dia_vfold,  
  param_info = rf_param)  
  
autoplot(rf_search, metric = "rmse") +  
  labs("Results of Grid Search for Two Tuning Parameters of a Random Forest")
```



# Model Selection

```
print(show_best(rf_search, "rmse", 9))
```

```
## # A tibble: 9 x 8
##   mtry degree .metric .estimator mean      n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
## 1     5     2 rmse    standard  0.124     3 0.00328 Preprocessor1_Model3
## 2     5     3 rmse    standard  0.125     3 0.00344 Preprocessor2_Model3
## 3     5     4 rmse    standard  0.126     3 0.00338 Preprocessor3_Model3
## 4     4     2 rmse    standard  0.126     3 0.00322 Preprocessor1_Model2
## 5     4     3 rmse    standard  0.127     3 0.00347 Preprocessor2_Model2
## 6     4     4 rmse    standard  0.127     3 0.00331 Preprocessor3_Model2
## 7     3     3 rmse    standard  0.131     3 0.00278 Preprocessor2_Model1
## 8     3     4 rmse    standard  0.132     3 0.00296 Preprocessor3_Model1
## 9     3     2 rmse    standard  0.132     3 0.00319 Preprocessor1_Model1
```

```
print(select_best(rf_search, metric = "rmse"))
```

```
## # A tibble: 1 x 3
##   mtry degree .config
##   <int> <int> <chr>
## 1     5     2 Preprocessor1_Model3
```

```
print(select_by_one_std_err(rf_search, mtry, degree, metric = "rmse"))
```

```
## # A tibble: 1 x 10
##   mtry degree .metric .estimator mean      n std_err .config      .best .bound
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>    <dbl> <dbl>
## 1     4     2 rmse    standard  0.126     3 0.00322 Preprocessor~ 0.124 0.127
```

# Best Model and Final Predictions

```
rf_param_final <- select_by_one_std_err(rf_search, mtry, degree, metric = "rmse")  
rf_wflow_final <- finalize_workflow(rf_wflow, rf_param_final)  
rf_wflow_final_fit <- fit(rf_wflow_final, data = dia_train)
```

- Want to use 'predict()' on data never seen before ('dia\_test');
- However, it does not work, because the outcome is modified in the recipe via 'step\_log()'.

# Best Model and Final Predictions

- Workaround:
  - ① Prepped recipe is extracted from the workflow;
  - ② This is used to 'bake()' the testing data;
  - ③ Use this baked data set together with extracted model for final predictions.

```
dia_rec3 <- extract_recipe(rf_wflow_final_fit)
rf_final_fit <- extract_fit_parsnip(rf_wflow_final_fit)

dia_test$.pred <- predict(rf_final_fit,
                          new_data = bake(dia_rec3, dia_test))$.pred
dia_test$logprice <- log(dia_test$price)

metrics(dia_test, truth = logprice, estimate = .pred)
```

```
## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 rmse    standard      0.113
## 2 rsq     standard      0.988
## 3 mae     standard      0.0846
```